

EDA

December 17, 2024

```
[31]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

# Read the CSV file
data_frame = pd.read_csv('../data/Customer_Churn.csv')
```

This is a detailed info about the features of the dataset, we can learn many things from this, for example, we get a sense of which are the nominal (Categorical) features (with object datatypes) and which are not. It also shows us how many records (rows) we have in this dataset, in this case, it's 3150.

```
[32]: # Basic information about the dataset
print("\n=== Dataset Info ===")
print(data_frame.info()) # Shows data types and missing values
```

```
=== Dataset Info ===
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3150 entries, 0 to 3149
Data columns (total 13 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   ID                                    3150 non-null   int64
 1   Call Failure                         3150 non-null   int64
 2   Complains                            3150 non-null   object
 3   Charge Amount                       3150 non-null   int64
 4   Freq. of use                        3150 non-null   int64
 5   Freq. of SMS                        3150 non-null   int64
 6   Distinct Called Numbers             3150 non-null   int64
 7   Age Group                           3150 non-null   int64
 8   Plan                                3150 non-null   object
 9   Status                              3150 non-null   object
10   Age                                 3150 non-null   int64
11   Customer Value                      3150 non-null   float64
12   Churn                              3150 non-null   object
dtypes: float64(1), int64(8), object(4)
```

memory usage: 320.0+ KB
None

```
[33]: # Summary statistics for all numeric columns
print("\n=== Summary Statistics ===")
print(data_frame.describe().round(2)) # Shows count, mean, std, min, 25%, 50%, 75%, max
```

```
=== Summary Statistics ===
```

	ID	Call Failure	Charge Amount	Freq. of use	Freq. of SMS	\
count	3150.00	3150.00	3150.00	3150.00	3150.00	
mean	1575.50	7.63	129.88	69.46	73.17	
std	909.47	7.26	102.79	57.41	112.24	
min	1.00	0.00	20.00	0.00	0.00	
25%	788.25	1.00	50.00	27.00	6.00	
50%	1575.50	6.00	100.00	54.00	21.00	
75%	2362.75	12.00	200.00	95.00	87.00	
max	3150.00	36.00	400.00	255.00	522.00	

	Distinct Called Numbers	Age Group	Age	Customer Value
count	3150.00	3150.00	3150.00	3150.00
mean	23.51	2.83	31.00	470.97
std	17.22	0.89	8.83	517.02
min	0.00	1.00	15.00	0.00
25%	10.00	2.00	25.00	113.80
50%	21.00	3.00	30.00	228.48
75%	34.00	3.00	30.00	788.39
max	97.00	5.00	55.00	2165.28

This is pretty helpful, as we learn there's no missing values in this specific dataset, which saves us the trouble of trying to handle them.

```
[34]: # Check for missing values
print("\n=== Missing Values ===")
print(data_frame.isnull().sum())
```

```
=== Missing Values ===
```

ID	0
Call Failure	0
Complains	0
Charge Amount	0
Freq. of use	0
Freq. of SMS	0
Distinct Called Numbers	0
Age Group	0
Plan	0
Status	0

```
Age                                0
Customer Value                     0
Churn                              0
dtype: int64
```

Sample of the dataset now that we understand the main characteristics of the dataset.

```
[35]: # Display first few rows
print("\n=== First Few Rows ===")
print(data_frame.head())
```

```
=== First Few Rows ===
   ID  Call Failure  Complains  Charge Amount  Freq. of use  Freq. of SMS  \
0    1             3         no           100           25           32
1    2             8         no           100           65            0
2    3             0         no           200            0            0
3    4            10         no           100           54          327
4    5            10         no           100           60            0

   Distinct Called Numbers  Age Group  Plan  Status  Age  \
0              11          3  pre-paid  active  30
1              13          2  pre-paid  active  25
2               0          2  pre-paid not-active  25
3              20          2  pre-paid  active  25
4              31          1  pre-paid  active  15

   Customer Value  Churn
0          193.120    no
1          194.400   yes
2           0.000   yes
3         1579.140   yes
4          227.865   yes
```

This helps us see how often does the customers churn, for example, notice that even though most of the customers don't complain (92%), it's not a good indicator that they're satisfied, as we see in the churn results, that around 84% of them do churn...we can see as well that most customers favour pre-paid plan(92%) over post-paid.

```
[36]: # Basic statistics for categorical columns
print("\n=== Categorical Columns Summary ===")
categorical_columns = data_frame.select_dtypes(include=['object']).columns
for col in categorical_columns:
    print(f"\nCounts for {col}:")
    print(data_frame[col].value_counts())
    print(f"\nPercentages for {col}:")
    counts = data_frame[col].value_counts(normalize=True) * 100
    for value, percentage in counts.items():
        print(f"{value}: {percentage:.2f}%")
```

=== Categorical Columns Summary ===

Counts for Complains:

Complains

no 2909

yes 241

Name: count, dtype: int64

Percentages for Complains:

no: 92.35%

yes: 7.65%

Counts for Plan:

Plan

pre-paid 2905

post-paid 245

Name: count, dtype: int64

Percentages for Plan:

pre-paid: 92.22%

post-paid: 7.78%

Counts for Status:

Status

active 2368

not-active 782

Name: count, dtype: int64

Percentages for Status:

active: 75.17%

not-active: 24.83%

Counts for Churn:

Churn

yes 2655

no 495

Name: count, dtype: int64

Percentages for Churn:

yes: 84.29%

no: 15.71%

A visualised distribution of the churn class label. (task 2)

```
[37]: # Create a figure with multiple subplots
fig, (ax1) = plt.subplots(1,figsize=(15, 5))
```


```
# Plot 1: Bar plot of churn distribution
sns.countplot(data=data_file, x='Churn', ax=ax1, palette=['#ff9999', '#66b3ff'])
ax1.set_title('Distribution of Customer Churn')
ax1.set_xlabel('Churn Status')
ax1.set_ylabel('Count')

# Add count labels on top of each bar
for i in ax1.containers:
    ax1.bar_label(i)
```

/tmp/ipykernel_119577/1728760192.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(data=data_file, x='Churn', ax=ax1, palette=['#ff9999',
'#66b3ff'])
```



Task 3: The Age group is the independent variable, presented on the x-axis. We can learn from this graph that people of Group ages 2 and 3, (ranges around 25-45) are the most people customers who are subscribed in the first place, and they're also the most customers to churn.

```
[43]: # Create a figure with appropriate size
plt.figure(figsize=(12, 6))

# Create a histogram showing churn distribution for each age group
# Using sns.histplot for better visualization of categorical data
sns.histplot(data=data_frame,
             x='Age Group',
             hue='Churn',
             multiple="dodge", # This makes bars appear side by side
             shrink=0.8)      # This adjusts the width of the bars
```

```

# Customize the plot
plt.title('Distribution of Churn by Age Group', fontsize=12, pad=15)
plt.xlabel('Age Group', fontsize=10)
plt.ylabel('Count of Customers', fontsize=10)

# Add a grid for better readability
plt.grid(axis='y', linestyle='--', alpha=0.7)

# Adjust layout to prevent label cutoff
plt.tight_layout()

# Show the plot
plt.show()

```

