

EDA

December 21, 2024

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

# Read the CSV file
data_frame = pd.read_csv('../data/Customer_Churn.csv')
```

```
/usr/lib/python3/dist-packages/scipy/__init__.py:146: UserWarning: A NumPy
version >=1.17.3 and <1.25.0 is required for this version of SciPy (detected
version 1.26.4
```

```
warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")
```

This is a detailed info about the features of the dataset, we can learn many things from this, for example, we get a sense of which are the nominal (Categorical) features (with object datatypes) and which are not. It also shows us how many records (rows) we have in this dataset, in this case, it's 3150.

```
[2]: # Basic information about the dataset
print("\n=== Dataset Info ===")
print(data_frame.info()) # Shows data types and missing values
```

```
=== Dataset Info ===
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 3150 entries, 0 to 3149
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
0	ID	3150 non-null	int64
1	Call Failure	3150 non-null	int64
2	Complains	3150 non-null	object
3	Charge Amount	3150 non-null	int64
4	Freq. of use	3150 non-null	int64
5	Freq. of SMS	3150 non-null	int64
6	Distinct Called Numbers	3150 non-null	int64
7	Age Group	3150 non-null	int64

```

8   Plan                3150 non-null  object
9   Status              3150 non-null  object
10  Age                 3150 non-null  int64
11  Customer Value      3150 non-null  float64
12  Churn               3150 non-null  object
dtypes: float64(1), int64(8), object(4)
memory usage: 320.0+ KB
None

```

```

[3]: # Summary statistics for all numeric columns
print("\n=== Summary Statistics ===")
print(data_frame.describe().round(2)) # Shows count, mean, std, min, 25%, 50%,
↳ 75%, max

```

```

=== Summary Statistics ===

```

	ID	Call Failure	Charge Amount	Freq. of use	Freq. of SMS \
count	3150.00	3150.00	3150.00	3150.00	3150.00
mean	1575.50	7.63	129.88	69.46	73.17
std	909.47	7.26	102.79	57.41	112.24
min	1.00	0.00	20.00	0.00	0.00
25%	788.25	1.00	50.00	27.00	6.00
50%	1575.50	6.00	100.00	54.00	21.00
75%	2362.75	12.00	200.00	95.00	87.00
max	3150.00	36.00	400.00	255.00	522.00

	Distinct Called Numbers	Age Group	Age	Customer Value
count	3150.00	3150.00	3150.00	3150.00
mean	23.51	2.83	31.00	470.97
std	17.22	0.89	8.83	517.02
min	0.00	1.00	15.00	0.00
25%	10.00	2.00	25.00	113.80
50%	21.00	3.00	30.00	228.48
75%	34.00	3.00	30.00	788.39
max	97.00	5.00	55.00	2165.28

This is pretty helpful, as we learn there's no missing values in this specific dataset, which saves us the trouble of trying to handle them.

```

[4]: # Check for missing values
print("\n=== Missing Values ===")
print(data_frame.isnull().sum())

```

```

=== Missing Values ===
ID                0
Call Failure      0
Complains         0
Charge Amount     0

```

```

Freq. of use          0
Freq. of SMS          0
Distinct Called Numbers 0
Age Group             0
Plan                  0
Status                0
Age                   0
Customer Value        0
Churn                  0
dtype: int64

```

Sample of the dataset now that we understand the main characteristics of the dataset.

```

[5]: # Display first few rows
print("\n=== First Few Rows ===")
print(data_frame.head())

```

=== First Few Rows ===

	ID	Call Failure	Complains	Charge Amount	Freq. of use	Freq. of SMS	\
0	1	3	no	100	25	32	
1	2	8	no	100	65	0	
2	3	0	no	200	0	0	
3	4	10	no	100	54	327	
4	5	10	no	100	60	0	

	Distinct Called Numbers	Age Group	Plan	Status	Age	\
0	11	3	pre-paid	active	30	
1	13	2	pre-paid	active	25	
2	0	2	pre-paid	not-active	25	
3	20	2	pre-paid	active	25	
4	31	1	pre-paid	active	15	

	Customer Value	Churn
0	193.120	no
1	194.400	yes
2	0.000	yes
3	1579.140	yes
4	227.865	yes

This helps us see how often does the customers churn, for example, notice that even though most of the customers don't complain (92%), it's not a good indicator that they're satisfied, as we see in the churn results, that around 84% of them do churn...we can see as well that most customers favour pre-paid plan(92%) over post-paid.

```

[6]: # Basic statistics for categorical columns
print("\n=== Categorical Columns Summary ===")
categorical_columns = data_frame.select_dtypes(include=['object']).columns
for col in categorical_columns:

```

```

print(f"\nCounts for {col}:")
print(data_frame[col].value_counts())
print(f"\nPercentages for {col}:")
counts = data_frame[col].value_counts(normalize=True) * 100
for value, percentage in counts.items():
    print(f"{value}: {percentage:.2f}%")

```

=== Categorical Columns Summary ===

Counts for Complains:

Complains

no 2909

yes 241

Name: count, dtype: int64

Percentages for Complains:

no: 92.35%

yes: 7.65%

Counts for Plan:

Plan

pre-paid 2905

post-paid 245

Name: count, dtype: int64

Percentages for Plan:

pre-paid: 92.22%

post-paid: 7.78%

Counts for Status:

Status

active 2368

not-active 782

Name: count, dtype: int64

Percentages for Status:

active: 75.17%

not-active: 24.83%

Counts for Churn:

Churn

yes 2655

no 495

Name: count, dtype: int64

Percentages for Churn:

yes: 84.29%
no: 15.71%

A visualised distribution of the churn class label. (task 2)

```
[7]: # Create a figure with multiple subplots
fig, (ax1) = plt.subplots(1, figsize=(15, 5))

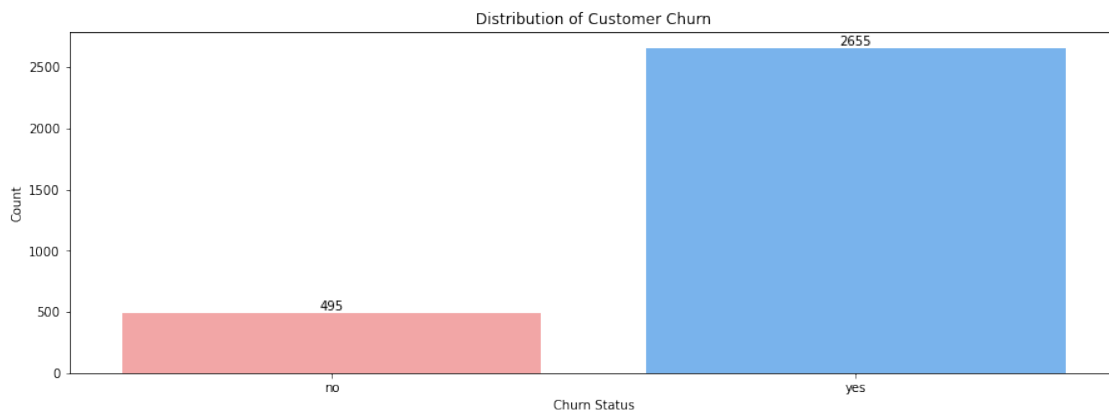
# Plot 1: Bar plot of churn distribution
sns.countplot(data=data_frame, x='Churn', ax=ax1, palette=['#ff9999', '#66b3ff'])
ax1.set_title('Distribution of Customer Churn')
ax1.set_xlabel('Churn Status')
ax1.set_ylabel('Count')

# Add count labels on top of each bar
for i in ax1.containers:
    ax1.bar_label(i)
```

/tmp/ipykernel_212184/852200220.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(data=data_frame, x='Churn', ax=ax1, palette=['#ff9999', '#66b3ff'])
```



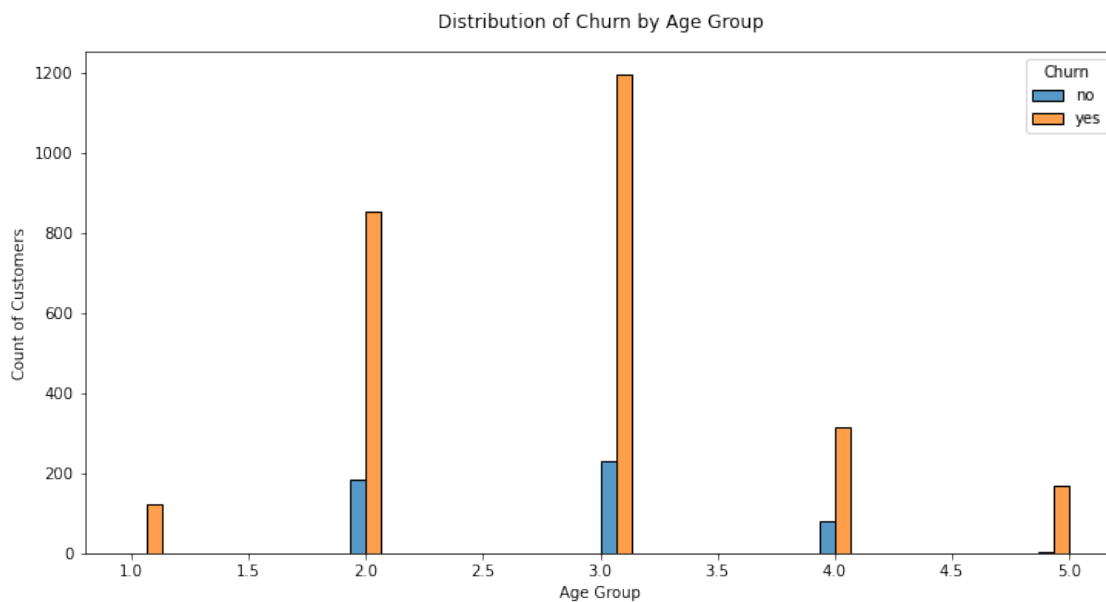
Task 3: The Age group is the independent variable, presented on the x-axis. We can learn from this graph that people of Group ages 2 and 3, (ranges around 25-45) are the most people customers who are subscribed in the first place, and they're also the most customers to churn.

```
[8]: # Create a figure with appropriate size
plt.figure(figsize=(12, 6))

# Create a histogram showing churn distribution for each age group
# Using sns.histplot for better visualization of categorical data
sns.histplot(data=data_frame,
             x='Age Group',
             hue='Churn',
             multiple="dodge" # This makes bars appear side by side
            )

# Customize the plot
plt.title('Distribution of Churn by Age Group', fontsize=12, pad=15)
plt.xlabel('Age Group')
plt.ylabel('Count of Customers')

# Show the plot
plt.show()
```



0.0.1 Task 4:

charge amount is the independent variable in this case.

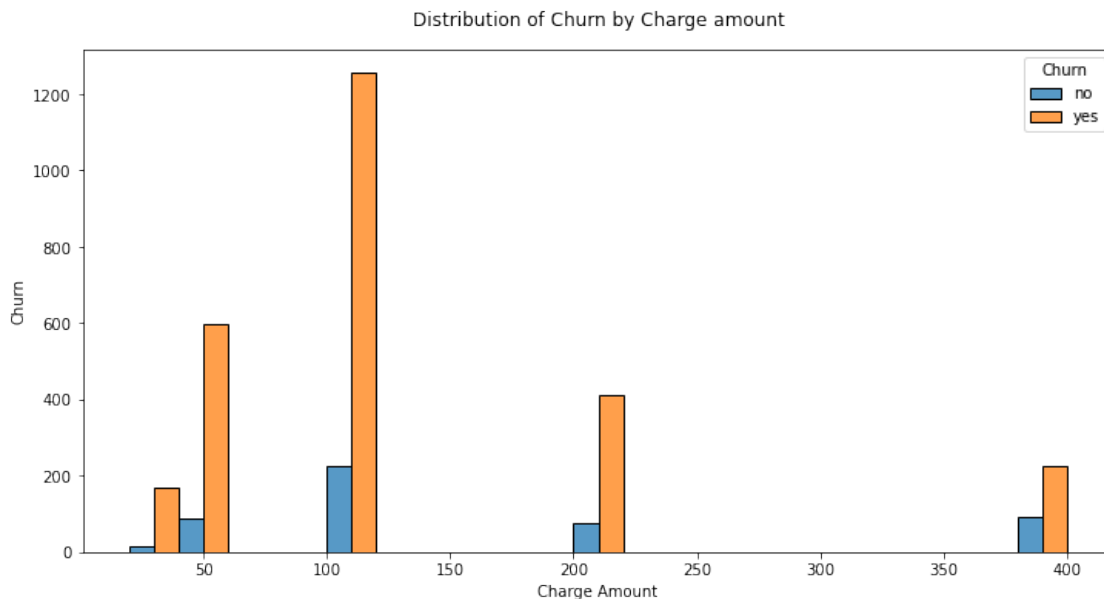
Conclusions It seems that the charge amount isn't the main cause for users to churn, because it doesn't matter which pricing they have, most users of this sub-group would churn eventually, we see charge amount with 100 and 50 are the most to churn, that's because they're the ones with the most users (Checkout charge amount analysis below)

```
[9]: plt.figure(figsize=(12,6))

# A histogram showing churn distribution among charge amount
sns.histplot(data=data_frame,
             x='Charge Amount',
             hue='Churn',
             multiple='dodge')

# Customizing the plot
plt.title("Distribution of Churn by Charge amount", fontsize=12, pad=15)
plt.xlabel("Charge Amount")
plt.ylabel("Churn")

plt.show()
```



0.0.2 Task 5:

Frequency of charge amounts and all the unique -amounts-, I thought there was something wrong with the graph above, given how separated the boxes, but now it's obvious, the Charge amounts are discrete values, which explains the discrete nature of the graph...

```
[10]: # Show unique charge amounts
print("Unique Charge Amounts in the dataset:")
print(sorted(data_frame['Charge Amount'].unique()))

# Show frequency of each charge amount
print("\nFrequency of each Charge Amount:")
```

```
print(data_frame['Charge Amount'].value_counts().sort_index())
```

Unique Charge Amounts in the dataset:

```
[20, 50, 100, 200, 400]
```

Frequency of each Charge Amount:

Charge Amount

```
20      184
```

```
50      683
```

```
100     1483
```

```
200      485
```

```
400      315
```

```
Name: count, dtype: int64
```

0.0.3 Task 6:

Tools and motivation I've decided to use a *heatmap* to visualise the correlations between features, because if I were to draw scatter plots for each two variables, I'll need (with 10 features) $n(n-1)/2 = 45$ plots! A heatmap can also help with avoiding *multicollinearity*, in the Assignment description: "Based on the correlation, you have to decide which features to stay for the learning stage and which can be deleted." This heatmap graph will be of great help on deciding which features to stay and which not, by avoiding *multicollinearity*.

First steps

1. First I need to map all -objects- types into a numeric representation, these are: "Complain, Status, Churn and Plan.

```
[11]: data_frame["Churn"] = data_frame["Churn"].map({'no': 0, 'yes': 1})
data_frame["Complains"] = data_frame["Complains"].map({'no': 0, 'yes': 1})
data_frame["Status"] = data_frame["Status"].map({'not-active': 0, 'active': 1})
# Get dummies here instead of map, because map creates an artificial numeric_
↳relationship
# between pre-paid and post-paid, which doesn't exist, unlike yes/no, active/
↳non-active
data_frame = pd.get_dummies(data_frame, columns=['Plan'])
```

2. Now, drawing the heatmap itself

Explanation:

- The values range from -1 (perfect negative correlation) to +1 (perfect positive correlation)
- 0 indicates no linear correlation

```
[12]: def plot_correlation_heatmap(data_frame):
      """
      Creates and displays a correlation heatmap for all numeric features in the_
      ↳dataset.
```


Parameters:

data_frame (pandas.DataFrame): The dataset to analyze, with ID column to be excluded

The function will:

- 1. Calculate correlations between all numeric columns (excluding ID)*
- 2. Create a masked triangular heatmap for better readability*
- 3. Display correlation values with 2 decimal places*

"""

Remove ID column and calculate correlation matrix

correlation_matrix = data_frame.drop('ID', axis=1).corr()

Set the visual style for better appearance

sns.set_theme(style="white")

Create a mask for the upper triangle

We do this because correlation matrices are symmetrical,

so we only need to show half to avoid redundancy

mask = np.zeros_like(correlation_matrix, dtype=bool)

mask[np.triu_indices_from(mask)] = True

Create the figure with a reasonable size

plt.figure(figsize=(11, 9))

Generate a blue-red color palette centered at 0

cmap = sns.diverging_palette(220, 10, as_cmap=True)

Create and customize the heatmap

sns.heatmap(correlation_matrix,

mask=mask,

cmap=cmap,

vmax=1,

center=0,

square=True,

linewidth=.5,

cbar_kws={'shrink': .5},

annot=True,

fmt='.2f')

Apply the triangular mask

Use our custom colormap

Set maximum correlation value

Center the colormap at 0

Make cells square

Add thin lines between cells

Customize the colorbar

Show correlation values

Format to 2 decimal places

Add title and adjust layout

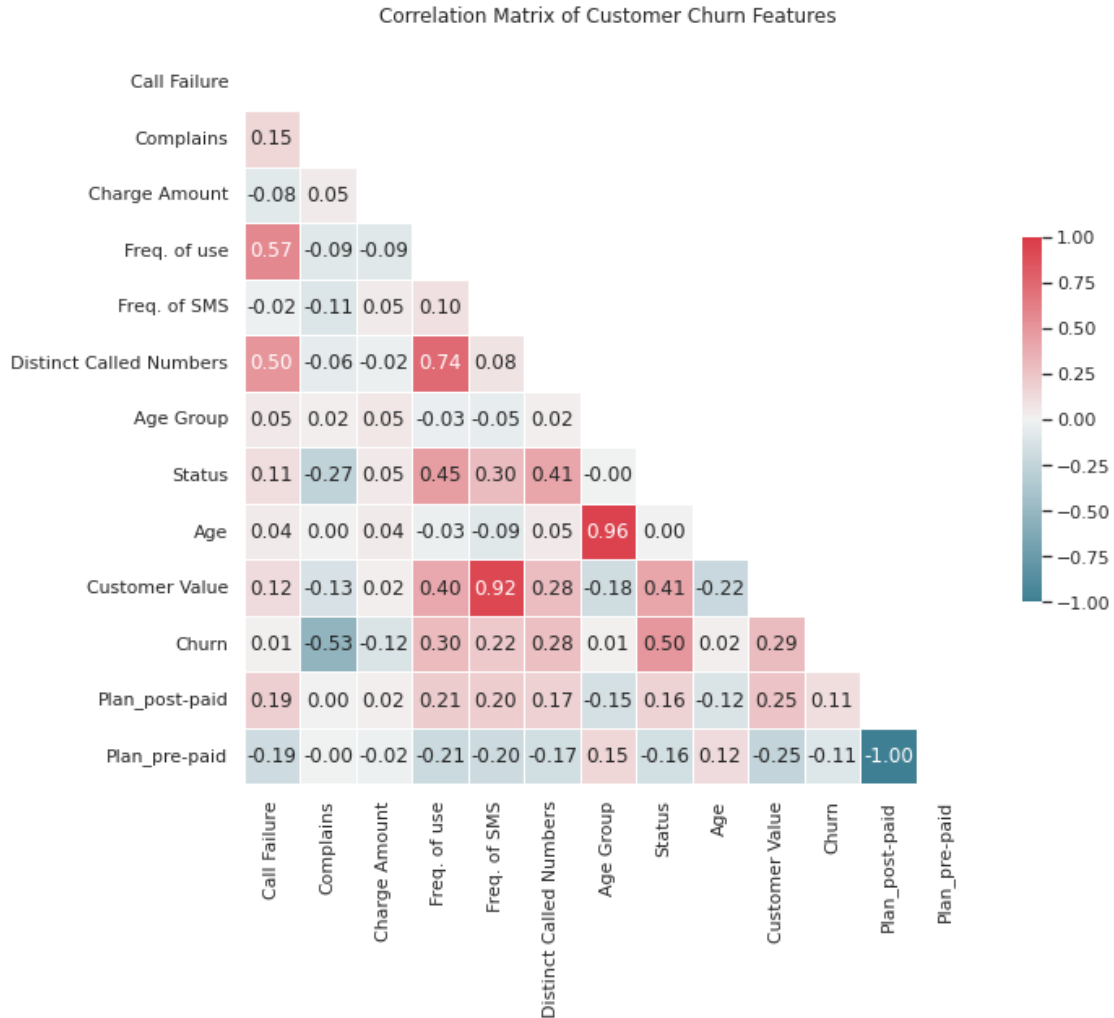
plt.title('Correlation Matrix of Customer Churn Features', pad=20)

plt.tight_layout()

Display the plot

plt.show()

plot_correlation_heatmap(data_frame)



0.0.4 Analysis of the heatmap plot:

1. Very high correlation between Customer value and Freq. of SMS, this requires to drop one of them before starting to train the model to avoid multicollinearity.
2. Pre-paid or Post-paid columns, one of them will be dropped, because if it's not pre-paid, then it can be concluded directly that it's postpaid.
3. For deciding between keeping Age or Age Group: if age group 2 (let's say ages 25-35) shows different churn behavior than age group 3 (ages 36-45), the exact transition point might actually be at age 32, not 35. By using the continuous Age variable, we allow our models to find these natural breakpoints in the data. That's why Age Group will be dropped, and Age continuous field will be kept.

```
[13]: # Dropping unwanted columns
      # set axis to 1 indicating that we're dropping columns, not rows (0)
```

```

data_frame = data_frame.drop(['ID', 'Age Group', 'Plan_post-paid', 'Freq. of
↳SMS'], axis=1)

# Verify the columns were dropped
print("Remaining columns:")
print(data_frame.columns.tolist())

# Re-print the heatmap
plot_correlation_heatmap(data_frame)

```

Remaining columns:

['Call Failure', 'Complains', 'Charge Amount', 'Freq. of use', 'Distinct Called Numbers', 'Status', 'Age', 'Customer Value', 'Churn', 'Plan_pre-paid']

```

-----
KeyError                                Traceback (most recent call last)
/tmp/ipykernel_212184/2985881702.py in <module>
      8
      9 # Re-print the heatmap
--> 10 plot_correlation_heatmap(data_frame)

/tmp/ipykernel_212184/1575683238.py in plot_correlation_heatmap(data_frame)
     12     """
     13     # Remove ID column and calculate correlation matrix
--> 14     correlation_matrix = data_frame.drop('ID', axis=1).corr()
     15
     16     # Set the visual style for better appearance

~/.local/lib/python3.10/site-packages/pandas/core/frame.py in drop(self, labels
↳axis, index, columns, level, inplace, errors)
    5566         weight 1.0      0.8
    5567         """
-> 5568         return super().drop(
    5569             labels=labels,
    5570             axis=axis,

~/.local/lib/python3.10/site-packages/pandas/core/generic.py in drop(self,
↳labels, axis, index, columns, level, inplace, errors)
    4783         for axis, labels in axes.items():
    4784             if labels is not None:
-> 4785                 obj = obj._drop_axis(labels, axis, level=level,
↳errors=errors)
    4786
    4787         if inplace:

~/.local/lib/python3.10/site-packages/pandas/core/generic.py in _drop_axis(self
↳labels, axis, level, errors, only_slice)

```

```

4825         new_axis = axis.drop(labels, level=level, errors=errors
4826     else:
-> 4827         new_axis = axis.drop(labels, errors=errors)
4828         indexer = axis.get_indexer(new_axis)
4829

~/local/lib/python3.10/site-packages/pandas/core/indexes/base.py in drop(self,
-> labels, errors)
7068         if mask.any():
7069             if errors != "ignore":
-> 7070                 raise KeyError(f"{labels[mask].tolist()} not found in_
-> axis")
7071             indexer = indexer[~mask]
7072         return self.delete(indexer)

KeyError: '['ID'] not found in axis"

```

0.0.5 Analysis of the new plot

1. Complains and Churn (-0.53): This is particularly interesting because it suggests that customers who formally complain are actually less likely to churn. This might indicate that customers who take the time to complain are more invested in the service and are giving the company a chance to address their concerns.
2. Frequency of use and Call Failure (0.57): This moderate positive correlation suggests that more frequent users naturally experience more call failures, which is expected given higher usage.
3. Frequency of use and Distinct Called Numbers (0.74): This remains the strongest positive correlation, indicating that customers who make more calls tend to contact more unique numbers.

0.0.6 Final Step, splitting the data

This is where sickit library comes in handy. The training and testing data set will be saved as a .npy files for quick processing during training and testing the models.

```

[15]: # This way, X has all the features that we might want to distribute the Churn_
-> feature across
X = data_frame.drop('Churn', axis=1)
Y = data_frame['Churn']

# Creating the split
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.3,
-> random_state=42)

# Save these splits for later use with training and testing the models
np.save('../data/X_train', X_train)

```

```
np.save('../data/X_test', X_train)
np.save('../data/Y_train', Y_train)
np.save('../data/Y_test', Y_test)
```